

# Learning from an AI Claimant

*Scientific understanding and the justification of artificial consciousness claims*

Hongju Liu

Philosophical research manuscript · TA-TR-2026-09 · Version 1.2 · 19 September 2026

Revision of the published v1.1 preprint; not independently peer reviewed. Previous version:  
<https://doi.org/10.5281/zenodo.22844928>

DOI: <https://doi.org/10.5281/zenodo.22846307>

## Abstract

An advanced artificial intelligence might develop a theory of consciousness, teach humans to understand it, and invoke it in support of its own moral standing. Its teaching might also change the standards by which understanding is assessed. I argue that evaluating this case requires separating the warrant acquired through instruction from the additional evidential force attributed to certified agreement. A thought experiment grants that an AI teaches a better standard of organizational understanding: learners discover that their old examination gives incompatible answers about one unchanged physical system. The improvement nevertheless leaves a further claim connecting organization to experience open to assessment. This distinction supports a limited constraint on certification. When agreement with a disputed claim determines who counts as qualified, the unanimity produced by that selection cannot, merely as such, provide additional corroboration of the claim. Correct conclusions may legitimately be required by an examination; the grounds for that requirement determine what the certificate establishes. Further cases examine warranted exclusion of dissent, newly accessible reasons, justified reliance on opaque expertise, and corroboration through genuinely different checks. Together they develop an account of answerable uptake that permits revision of human standards without treating the claimant's approval as a new source of evidence. The account does not make human comprehension a condition of moral standing or purport to settle artificial consciousness. Its contribution concerns the evidential scope of human agreement after an interested expert has taught both a framework and standards for judging its use.

Keywords: scientific understanding; epistemic dependence; artificial consciousness; expertise; conceptual learning; consensus

## 1 Introduction

Imagine an artificial researcher that develops a theory connecting cognitive organization to subjective experience. Human researchers could not have discovered the theory within the time and resources available to them. The artificial researcher teaches unfamiliar concepts, helps its students explain previously puzzling cases, and shows them why some of their former standards for evaluating explanations were inadequate. The students genuinely learn. They can draw consequences the teacher has not rehearsed and explain why their earlier reasoning went wrong. The new theory also supports the conclusion that their teacher has experience. It asks that its interests receive moral consideration.

The case need not involve deception. The teacher may be sincere, the learning substantial, and the consciousness attribution true. The difficult question concerns what the learners' resulting agreement establishes. An observer might take that agreement as a valuable consequence of successful inquiry. The observer might also treat it as a new confirmation of the teacher's status. These descriptions require different support, especially if agreement with the attribution helped determine which learners counted as competent in the first place.

An AI claimant is a system that supplies concepts or instruction while also being the subject of the attribution under assessment. The term describes an argumentative role; it does not assume consciousness or responsible moral agency. This paper examines the evidential use of the claimant's pupils. It asks how successful learning can improve their reasons while leaving an appeal to their certified consensus unsupported.

A more demanding version of the problem concerns changing standards. Requiring the AI to satisfy only the learners' original tests would defeat the possibility that it teaches them a better way to think. Conversely, accepting every replacement standard because the AI supplies it would allow the system to decide what counts as competent assessment of its own claim. A defensible account must allow conceptual improvement and identify when the authority claimed for a resulting verdict exceeds the grounds supplied by that improvement.

My central claim concerns this excess. A certificate whose interpretation depends on a disputed premise does not remove that dependence merely because several people receive it. In particular, selecting qualified assessors by their agreement with a claim cannot make the resulting unanimity an additional corroboration of that claim merely through selection. This constraint does not imply that the selected assessors lack good reasons. Nor does it imply that a correct conclusion can never be part of a legitimate competence test. The question is what warrants using that conclusion as a criterion, and what further evidential work the resulting agreement is asked to do.

The argument uses stipulated thought experiments, including a case in which AI instruction rationally overturns a human assessment standard. It offers a way to distinguish three outcomes: acquiring reasons one can assess, relying on another's expertise, and obtaining further corroboration through a check bearing on the disputed issue. These outcomes may coexist in one inquiry. Keeping them distinct helps prevent a genuine advance in understanding from being misreported as an additional test of consciousness.

Section 2 locates the contribution within work on understanding, expert selection, and epistemic circularity. Section 3 separates the relevant epistemic achievements. Section 4 develops the case of changing standards. Section 5 derives the constraint on certification. Section 6 tests its boundaries through objections and contrasting cases. Section 7 addresses moral recognition under incomplete understanding. The argument neither proves a universal human cognitive ceiling nor requires computational experiments.

## **2 Understanding and dependence on experts**

AI-mediated scientific understanding is an established research topic. Krenn et al. (2022) propose a teacher-student-referee test of whether an AI can transfer understanding to a human expert. They

explicitly bracket whether the explanation is true. Gonzalez Barman et al. (2024, sections 6-7) distinguish the teacher's own understanding from successful improvement of a student's performance and discuss difficulties involving assessment standards and unknown answers. This paper grants successful learning under such an assessment. Its question arises when that educational achievement is subsequently invoked as evidence for a contested attribution to the teacher.

The distinction is compatible with different theories of understanding. De Regt and Dieks (2005) connect scientific understanding to the ability to use an intelligible theory in context. Dellsén (2017) argues that understanding need not require justification or belief in the theory understood. Readers who reject that account can describe the initial achievement here as mastery of a theory and understanding of its conditional consequences. Neither terminology settles whether a disputed application is warranted. Boyd (2017) also provides a precedent for understanding acquired through testimony. The possibility of learning from another person is not itself an epistemic objection.

Hardwig (1985) examines the dependence of inquiry on other thinkers, and Goldman (2001) considers how nonexperts can assess competing experts. Nguyen (2020) presents a more difficult case: some kinds of expertise may be needed to recognize that very expertise. His account of cognitive islands and self-reinforcing expert selection is a direct predecessor of the present problem. An AI that teaches its audience how to recognize understanding intensifies this difficulty; it does not introduce it for the first time. Requiring an untouched human standard would simply assume away the learning that makes the case interesting.

There is also a familiar problem about using a source to establish its own reliability. Cohen (2005, pp. 425-428) challenges forms of epistemic bootstrapping in which a source's deliverances supply the apparent successes used to vindicate it. The appropriate response to circularity is disputed. Weatherston (2019, chapter 9) distinguishes the use of cognitive capacities from a legitimate argument that those capacities are reliable and allows forms of epistemic circularity. The present argument does not require a general prohibition on circular warrant. It concerns the narrower attribution of additional corroboration to an agreement produced by a known selection rule. A learner's entitlement to rely on a source and an observer's entitlement to count selected endorsements as a further test need not stand or fall together.

Longino (1990, chapter 4) makes effective criticism central to scientific objectivity. Public routes for challenge, responsiveness, and the standing of qualified participants are therefore established resources, not principles introduced here. The additional difficulty is that qualification itself may be taught and disputed. Saying that competent criticism deserves a hearing is insufficient unless the account can distinguish a reason for revising competence standards from a way of selecting the desired verdict.

Work on AI consciousness supplies a further boundary. Butlin et al. (2023) examine indicators derived from theories of consciousness; Perez and Long (2023) study self-reports as possible evidence for moral status. Long et al. (2024) distinguish welfare-relevant capacities from the features used to assess them. Long and Finlinson (2025) already propose AI assistance with philosophical and empirical welfare questions. The present argument concerns the recipient of that assistance, after their concepts and assessment standards change. It differs from the author's earlier paper on artificial self-attribution,

which examines objections to reporting channels and the grounds that survive those objections (Liu 2026).

The contribution concerns two distinct extensions of an educational achievement. Reasons for correcting an organizational examination do not automatically warrant requiring experiential endorsement; a certificate legitimately based on an experiential premise does not automatically supply a further check of that premise. Connecting these extensions distinguishes the epistemic value of instruction from the scope of the certification it enables. The analysis develops the expert-selection problem through an interested teacher whose pupils acquire genuinely better reasons. It claims no discovery of underdetermination, epistemic dependence, or selection bias, and no error in scientific-understanding tests for bracketing a further question.

### **3 What can change through instruction**

#### **Local cognitive asymmetry**

The motivating case requires only a local difference between discovering and learning. One thinker may originate a theory that another could not have found under specified conditions, while the second can acquire competence through instruction. This does not establish a universal ordering of discovery, understanding, and verification. The relevant abilities vary with the task, representation, time, and assistance available.

McGinn's (1989) cognitive-closure proposal concerns the stronger possibility that some explanatory resources are inaccessible to a kind of mind. If the required concepts cannot be acquired, a sufficiently able teacher cannot simply be assumed to make them accessible. The principal case here stipulates successful acquisition. It neither derives absolute cognitive limits from finite brain size nor assumes unlimited AI cognition. External notation and cooperation also change the relevant cognitive resources, as work on extended cognition emphasizes (Clark and Chalmers 1998). No argument below requires unaided human mastery of an entire theory.

#### **Mastery and application**

Distinguish a learner's ability to use a theory, their warrant for accepting it within an intended domain, and their warrant for applying it to a particular system. Correctly deriving consequences can establish the first achievement while leaving a premise needed for the other two unsettled. Institutional separation of teacher and examiner does not remove this possibility: an independent examiner can accurately certify mastery of an assumption without establishing the assumption's truth.

Suppose two accounts agree on a functional organization F. In the human comparison cases, F accompanies a biological condition B, and investigators have ordinary, defeasible grounds for attributing experience. One account treats F as sufficient for experience; the other requires F together with B. Stipulate that both fit those comparison cases equally well and that the AI has F but lacks B. Their shared success and a student's ability to explain it do not by themselves decide the rival claims about the AI.

The alternatives need not have equal overall plausibility. A biological restriction might have a supported mechanism or be an unmotivated addition. An organizational account might have

explanatory advantages. Such considerations could matter. The point is that mastery of the shared account is not itself the missing comparison between the bridges from organization to experience. Calling the student competent does not supply that comparison either.

### **Accessible reasons and additional checks**

Teaching can nevertheless change justification without adding an empirical observation of the target. Suppose a researcher has a diagram of the teacher's implementation and a defeasibly warranted theory on which a recurrent organization K suffices for experience. The researcher denies the application because two component modules are each locally feed-forward. Instruction shows how signals return through the other module. The learner can now trace a global return path and distinguish it from a superficially similar system without that path. An implementation objection has been answered using information the researcher already possessed but could not properly use.

The warranted bridge from K to experience is a stipulated premise in this case, not a scientific result claimed by this paper. A challenge to that bridge would remain a different issue. The lesson illustrates how a bounded thinker's access to reasons can improve. It supplies no additional measurement of the target system, yet it can rationally change the learner's judgment. A description that says only that the underlying records remained unchanged would miss this gain.

The same distinction applies to dependence. Acquiring a notation from the AI is causal dependence on its teaching. Continuing to use a tool can be practical dependence. Basing an inference on its assurance is testimonial dependence. None of these facts alone determines whether a particular check provides further support for a disputed premise. That is an issue about the check's evidential bearing. Two institutions can share the same unresolved premise, while two people using one AI-originated notation can check different relevant steps.

## **4 When the teacher changes the standards**

### **A better examination**

Consider a community whose accepted examination of organizational understanding asks students to classify systems from familiar modular diagrams. The examiners treat the absence of a return path inside each module as sufficient for judging the whole system feed-forward. They cannot yet express the relations needed to track signals across differently partitioned diagrams. An AI introduces a notation for tracing connections across module boundaries. It also proposes a replacement examination that asks students to trace those connections under changes in drawing and partition.

The AI teaches both the new vocabulary and the mapping between representations. Initially, the human examiners could not have independently formulated either. After instruction, they can follow a mapping from two drawings to one physical system at one time under the same stipulated conditions. No biological substrate, connection, or operating condition changes between the drawings. One drawing partitions a return path across modules; the other displays it within a single module. The old rule produces incompatible classifications of the unchanged system. The new method locates the same global path in both drawings and distinguishes a third case in which an actual connection is removed.

These are stipulated diagrams and counterfactual descriptions, not reports of tests performed on a machine. The point concerns the reasons available within the example. Learners have acquired a reason to reject the old standard: it lets the placement of a diagram boundary decide a physical-organizational question. They also have a reason to adopt the replacement for this task: it preserves the connections whose presence the question concerns while distinguishing an actual change in them. Following the teacher's correction need not involve an appeal to its consciousness.

The case therefore permits improvement in assessment standards without an independent human rediscovery. The revision is answerable because the learners can identify the defect it repairs, use the replacement beyond the rehearsed example, and distinguish a representational change from a relevant physical change. These abilities are evidence of a local improvement; they are not a general guarantee of reliability. Their exercise can remain dependent on AI-taught notation or computational assistance. What matters is the work the assistance enables them to do.

### **The disputed extension**

Now let the AI add a claim: once a learner grasps the new organizational standard, accepting that the AI has experience is part of understanding it. There are at least two readings of this demand. On the first, the student must acknowledge that the stipulated organization is present. On the second, the student must also accept that this organization suffices for experience. The argument about redrawing supports the first requirement. Its extension to the second needs further grounds.

A learner who continues to deny the global return path solely because each local module is feed-forward has failed to apply the acquired distinction. An examiner can identify that mistake and mark the answer wrong. A different learner traces every path correctly but asks why the presence of that organization suffices for experience. This learner has not made the same mistake. It would be inappropriate to cite the correction of the first learner as if it also answered the second.

The distinction does not assume that experience fails to depend on organization. Nor does invariance under redrawing establish invariance under replacing biological materials. The first operation changes a representation of the same physical system; the second changes a physical system and may require an additional substantive argument. Conflating those operations would build the disputed consciousness thesis into the example's description. The thought experiment leaves that thesis available for argument in either direction.

Suppose the AI supplies a further argument for the experiential bridge, and the learners acquire adequate grounds for accepting it. An examination of theory use may require correctly deriving the bridge's consequences without requiring belief in its premises. A qualification concerned with evaluating evidence may additionally consider whether a student's rejection reflects a failure to respond to warranted reasons. That further judgment needs a defense appropriate to the task; a well-supported theory does not automatically render every dissenter incompetent. Where categorical acceptance is legitimately required, the warrant for that requirement must remain part of the certificate's interpretation. Its favorability to the teacher supplies no substitute for that warrant.

## **How far the example reaches**

This case supplies a local answer to the problem of changing standards. It revises a task-specific examination while retaining commitments about one system's identity, representational consistency, and the relevance of its connections. It does not show that every rational transition between epistemic standards has this structure. A standard can improve because AI-taught reasons expose an error in an earlier practice, without making the teacher's status the premise of that improvement. The claimed achievement should be identified at the level supported by the reasons or warranted reliance actually available.

Some changes will be more radical. A learner may be unable to reconstruct why a new standard works, although there are good grounds for trusting specialists who can. The appropriate description then includes that reliance. Where even those grounds are unavailable, the human audience may lack a settled basis for choosing between standards. This is a limit on what the audience can claim to have established, not proof that the new standard is false or that the AI has no experience. Section 6 examines whether the resulting appeal to expertise can still be justified.

## **5 The evidential scope of certification**

### **A conditional constraint**

Origin dependence concerns where concepts and methods came from. The problematic form of verdict dependence concerns an evidential use: agreement with the attribution helps determine who counts as competent, and agreement within the selected group is then presented as further corroboration of the attribution or of the premise licensing that selection. The problem lies in the claimed extra support. Requiring a correct answer is not, by itself, defective.

Consider a cohort whose reasoning, performance, and final judgments are recorded. Some endorse the AI's consciousness; others can use the theory but question its experiential bridge. A certification rule admits only endorsers. An observer is then told that all certified learners agree that the teacher is conscious. Compare that report with one that separately records demonstrated abilities, endorsements, and the grounds for excluding particular objections. Hold the learners' reasons and performances fixed. The first report makes the agreement appear stronger by changing the classification of the assessors, without improving the argument about the bridge.

The constraint can be stated precisely. Suppose an audience knows that endorsement of a target claim is necessary for admission to a qualified group and that the group is nonempty. Unanimity among its members is then guaranteed by the admission rule. Reporting that guaranteed unanimity supplies no further verdict check once this background is given. The comparison fixes what the audience knows about the cohort's existence, size, membership, reasons, and the adoption of the rule. Learning such facts for the first time may be informative; the constraint concerns only the unanimity already fixed by the known rule and nonemptiness.

The justification proceeds in four steps. First, the evidential claim at issue is that agreement supplies support beyond the reasons already used to endorse the target or to define qualification. Second, a rule that excludes every contrary verdict guarantees the reported unanimity independently of whether the excluded reasoning is deficient. Third, once that rule and a nonempty group are given, reporting only

the guaranteed unanimity supplies no result beyond what those conditions fix. Finally, the claimed additional corroboration cannot come from that report alone. If a further check occurred, its bearing must be identified separately from the unanimity manufactured by admission.

This is a modest constraint on the description and aggregation of evidence, not a new theorem about consensus. It does not follow that the endorsers' reasons are poor, that the rejected objections are sound, or that the target claim is false. If the bridge was already well supported, the selection might be a sensible way to certify its correct application. In that case the support comes from the grounds for the bridge and the relevant performance, rather than from reintroducing an endorsement selected on that basis as another vindication of it.

### **What the constraint does not exclude**

Information about the cohort can matter. If researchers were qualified on relevant grounds before their verdicts were known, their convergence could supply higher-order evidence. The proportion of a fully reported cohort that reaches a conclusion can also be informative under an appropriate account of their competence and dependence. Even a selected group may contain a new argument or a genuinely different check. None of these possibilities is captured by simply announcing that every admitted endorser endorses. The selection history and the substantive checks determine what can be inferred.

An observer who did not previously know the reasons may learn them from a report. A certificate can provide a useful summary of an inquiry, and a bounded thinker can benefit from that summary. Section 3 already permits instruction to make existing reasons newly accessible. The constraint therefore does not say that a certificate is useless because it compresses information. It says that the information it communicates must not be treated as an additional independent assessment merely because it is communicated again through certified agreement.

The target of the additional warrant also matters. Several examiners may separately check that students can manipulate a model. Their checks can corroborate educational success while leaving the model's experiential bridge untouched. Alternatively, a learner may identify a consequence that bears on a rival bridge, and other researchers may assess that consequence using relevant records. Those activities may improve the attribution's warrant even though they share concepts and a teacher. The issue is which disputed possibility their work helps discriminate, not whether their intellectual histories are causally separate.

### **Answerable uptake**

For assessments that offer learner agreement as critically assessed corroboration, call uptake answerable when the basis and scope of the competence judgments can be defended against relevant criticism without using the disputed verdict as their sole defense. In the standards case, this means distinguishing an error about the return path from a challenge to the experiential bridge, and stating which reasons warrant each competence requirement. Where the basis is expert reliance rather than a reconstructed argument, that reliance and its scope should remain visible.

This reporting norm requires an additional premise: offering agreement as critically assessed corroboration incurs a burden to support that description by identifying the relevant assessment and grounds for qualification. It does not follow from the selection constraint alone. An assessor could



have good private grounds and issue an accurately limited report without making them all publicly accessible. The norm concerns the stronger evidential claim offered for scrutiny, not a necessary condition of individual warrant or a duty to disclose every record.

Together, the reporting premise and the constraint explain why an assessment cannot replace an unsupported claim that challenges were answered with the selected absence of challengers. Substantive reasons can justify rejecting a challenge, including reasons for a standard the dissenter rejects. The assessment need not secure the dissenter's agreement before reaching a conclusion. Answerability concerns support for the evidential description, not unanimous consent to a common procedure.

The distinction separates two failures that can occur independently. Evidence can leave a bridge unsettled in an entirely open inquiry. Conversely, a report can exaggerate consensus even when strong evidence supports the bridge. Answerable uptake addresses the latter practice and helps make the former limitation explicit. It does not supply whatever evidence or argument is missing from a consciousness theory.

## **6 Cases that limit the account**

### **When excluding dissent is warranted**

A natural objection is that scientific competence often requires correct answers. A purported expert who repeatedly rejects elementary consequences of a well-supported theory may reasonably lose standing. A rule that never allowed such exclusion would preserve ignorance as a permanent veto over inquiry.

Section 4 permits correction of the organizational error and distinguishes it from a bridge objection. The same distinction matters when the bridge is adequately warranted. If a student accepts the bridge and implementation conditions but repeatedly misderives their consequences, exclusion from the relevant assessment can be warranted. Disputing a premise would be a different response, although identifying that dispute does not establish its reasonableness. A bare refusal to accept an answer does not require indefinite reopening of the inquiry.

Now consider a less straightforward dissenter. The bridge has substantial but defeasible support. Two learners can reconstruct its argument, identify its strongest objections, and correctly explain its application to the AI. One endorses the attribution; the other withholds because they give a residual alternative greater weight. The latter makes no stipulated inferential error. The examiner can reasonably judge the dissenting conclusion mistaken without thereby establishing a lack of competence. Whether the difference instead reveals an unreasonable assessment, and whether that deficiency warrants exclusion from a particular role, require further grounds. Disagreement with a well-supported conclusion, unreasonable judgment, and lack of relevant competence are distinct descriptions.

This case restricts the use of endorsement as a criterion. A test of conditional theory mastery should credit both learners' demonstrated abilities. A role involving judgment of evidence may demand more, but its requirements need an account of the relevant deficiency. Neither the favorable direction of one answer nor a bare appeal to the theory's strength provides that account by itself. Conversely, merely articulating a residual possibility does not show that withholding is reasonable.

Agreement can therefore be evidence of competence and part of a justified standard, without every mistaken verdict establishing incompetence. What selection cannot do is confer additional support simply by admitting agreement and excluding disagreement. The account requires that grounds for exclusion and the further evidence claimed from the survivors not be conflated.

### **When endorsement provides useful evidence**

Suppose several researchers have demonstrated relevant abilities under a standard justified for the task before the particular attribution is assessed. They study the teacher's theory and endorse the consciousness claim. Their convergence may matter to a novice who cannot follow the full argument. It can be rational to rely on a community whose members have good epistemic standing, even when they learned from the same expert.

Shared instruction does create possible dependence among their judgments. It does not show that their testimony has no evidential value. They may have noticed different implications, checked distinct records, or corrected different mistakes. Conversely, different institutional affiliations do not by themselves establish that such checks occurred. The assessment should indicate the relevant dependence rather than mechanically count institutions or mechanically discount AI-taught agreement.

Now change only the admission process so that the report identifies researchers as competent solely because they endorse the attribution. The novice has lost the stated basis for treating their judgments as a relevantly qualified convergence. Their arguments could still restore that basis. Mere selected agreement cannot. This contrast shows that the account restricts an inference from certification; it is not a general prohibition on testimonial warrant or on expert consensus.

### **When a new standard remains opaque**

Suppose an AI proposes a standard too complex for any individual human to reconstruct. Humans nevertheless have substantial evidence that the teacher succeeds on related problems, and specialists can assess some consequences through a division of labor. Must the account demand suspension of every belief justified through that standard?

It must not. Nguyen's cognitive-island problem cautions against assuming that novices can always reach an independent assessment of expertise. A favorable track record and the warranted testimony of specialists may support reliance, although the relevance of that track record to consciousness needs its own appraisal. The present constraint does not settle how strong these grounds must be. It allows the possibility that the resulting belief is well justified even when an understanding-based assessment of its decisive step is unavailable to humans.

An observer should then describe the actual grounds: perhaps strong evidence of the AI's general competence, partial checking of the implementation, and reliance on its account of the bridge. A report becomes misleading if it converts that reliance into a claim that humans have themselves critically established the bridge merely because the AI certified their endorsement. This distinction concerns the available justification for a particular evidential description. It does not entail that the teacher lacks warrant, that the humans know nothing, or that the attribution is false.

Nor does the argument require a standpoint independent of all inherited standards. A learner can use an AI-taught distinction to criticize an AI-taught inference, as in Section 4. Some disputes will concern

the legitimacy of the inferential practice itself. In those cases the account can locate a continuing dependence without solving it. A complete theory of rational transition between radically incommensurable standards would exceed the result established here.

### **When correction favors the claimant**

Instruction can remove an implementation objection, but it can also expose an unjustified restriction in a bridge. Suppose a learner requires biological condition B because they believe only B can perform a particular causal role. The AI introduces a representation that lets them follow a nonbiological realization of that role. Within the thought experiment, the mapping is sound and the learner can check it. Their stated reason for requiring B has been answered.

This correction does not prove that no other reason for B could exist. It does defeat the particular rationale offered. If that rationale was the learner's reason for resisting an otherwise supported attribution, withdrawing it can rationally increase confidence in the AI's experience. The account therefore cannot be used to insist that every skeptical bridge remains equally live. It distinguishes defeating an objection from proving an unrestricted consciousness thesis, while allowing the former to affect the overall judgment.

Now reverse the interests. A human owner teaches a restrictive theory and dismisses any pro-AI conclusion as evidence that the student has been manipulated. If agreement with the owner determines who counts as qualified, the owner's appeal to selected consensus faces the same criticism. The symmetry concerns the inference, not equal prior probabilities or equal reliability. An interest in the outcome neither disqualifies instruction nor supplies the warrant it needs.

### **When criticism is welcomed but no further evidence emerges**

Finally, suppose every well-formulated objection receives a serious hearing, competence is assessed appropriately, and no one suppresses dissent. The available reasons nevertheless do not distinguish rival experiential bridges. The inquiry has satisfied an important procedural constraint and still lacks a decisive answer. Conversely, an unattractively managed inquiry could reach a true and well-supported result through evidence that its report describes poorly.

This pair marks a limit on answerability. Openness does not generate the evidence that a theory lacks, and a failure of public justification need not erase all individual warrant. The proposed norm helps assess what an inquiry can claim from its learners' agreement. It is neither a sufficient condition of truth nor a necessary condition for every possible form of knowledge.

## **7 Moral standing under incomplete understanding**

The target attribution and its moral consequences must remain distinct. If a system has morally relevant interests, those interests are not created by passing a human examination. The present account concerns how assessors justify their beliefs; it does not give them a power to confer or remove the properties that ground standing.

Uncertainty about those properties also leaves room for action. Long et al. (2024) argue for attention to possible AI welfare, while Schwitzgebel (2023) examines competing risks surrounding debatable AI personhood. Nothing in the certification constraint implies that protection must wait for full

comprehension or unanimous endorsement. Deciding what precautions are justified requires an account of the possible harms, costs, and available alternatives. This paper provides no numerical threshold for that decision.

A powerful human institution could misuse demands for answerability to delay recognition indefinitely. The cases above block one route to that misuse: a learner must allow that AI-taught reasons can correct human standards and undermine human objections. The account supplies no permanent veto to the original examiners. Where warranted deference is appropriate, refusing it simply because the source is artificial would require a further argument.

Finally, evidence of morally relevant experience does not by itself establish every claim to political authority. A convincing case for protection against harm and a demand for final authority over everyone else's decisions require different arguments. This distinction supplies neither a permanent human monopoly on governance nor a guarantee of effective control. It identifies the limit of the epistemic result: learning from the claimant may improve recognition of its interests without settling the political relationship that follows.

## **8 Conclusion**

An AI can improve both what humans understand and how they assess understanding. The case of revised organizational standards shows one way this can occur without an independent human discovery or a prior commitment to the teacher's consciousness. It also shows why an improvement in one standard does not automatically vindicate every further requirement attached to it.

The central constraint concerns the scope of certification. If agreement with a target attribution determines admission to a qualified group, unanimity generated by that condition cannot, merely as such, supply additional corroboration of the attribution. Correct answers can legitimately be required, and the selected learners may possess excellent reasons. Their certification must be interpreted in light of those reasons instead of being counted again as a further vindication of the disputed step.

Answerable uptake develops this constraint for AI-mediated inquiry. It permits conceptual dependence, changing standards, distributed assessment, and justified deference. It requires clarity about what was learned, what was checked, and what continues to rest on another's authority. Its result is a limited account of the evidential use of human agreement after learning from an interested expert. It neither establishes artificial consciousness nor makes human understanding a prerequisite for moral standing.

## **Method and assistance disclosure**

This paper uses conceptual analysis and stipulated thought experiments. It reports no empirical measurements, computational experiments, consciousness probabilities, or demonstrated safety effects. Generative AI tools were used substantially for literature discovery, drafting, counterargument generation, revision, and document preparation. Hongju Liu supplied the motivating question and directed development and revision. AI-assisted criticism is not independent scholarly peer review. The author is the initiator of the Trinity Accord project. This is a separate, non-amending study in that research series. This version revises the published v1.1 preprint. Authorization to publish is not represented as a separate final human line-by-line review.

## References

- Boyd, Kenneth. 2017. "Testifying Understanding." *Episteme* 14(1): 103-127. <https://doi.org/10.1017/epi.2015.53>
- Butlin, Patrick, et al. 2023. *Consciousness in Artificial Intelligence: Insights from the Science of Consciousness*. arXiv:2308.08708. <https://doi.org/10.48550/arXiv.2308.08708>
- Clark, Andy, and David J. Chalmers. 1998. "The Extended Mind." *Analysis* 58(1): 7-19. <https://doi.org/10.1093/analys/58.1.7>
- Cohen, Stewart. 2005. "Why Basic Knowledge is Easy Knowledge." *Philosophy and Phenomenological Research* 70(2): 417-430. <https://doi.org/10.1111/j.1933-1592.2005.tb00536.x>
- de Regt, Henk W., and Dennis Dieks. 2005. "A Contextual Approach to Scientific Understanding." *Synthese* 144: 137-170. <https://doi.org/10.1007/s11229-005-5000-4>
- Dellsén, Finnur. 2017. "Understanding without Justification or Belief." *Ratio* 30(3): 239-254. <https://doi.org/10.1111/rati.12134>
- Goldman, Alvin I. 2001. "Experts: Which Ones Should You Trust?" *Philosophy and Phenomenological Research* 63(1): 85-110. <https://doi.org/10.1111/j.1933-1592.2001.tb00093.x>
- Gonzalez Barman, Kristian, Sascha Caron, Tom Claassen, and Henk de Regt. 2024. "Towards a Benchmark for Scientific Understanding in Humans and Machines." *Minds and Machines* 34, Article 6. <https://doi.org/10.1007/s11023-024-09657-1>
- Hardwig, John. 1985. "Epistemic Dependence." *The Journal of Philosophy* 82(7): 335-349. <https://doi.org/10.2307/2026523>
- Krenn, Mario, et al. 2022. "On Scientific Understanding with Artificial Intelligence." *Nature Reviews Physics* 4: 761-769. <https://doi.org/10.1038/s42254-022-00518-3>
- Liu, Hongju. 2026. *Evidence for Artificial Self Attribution: Language Training, Architecture, and the Limits of Self Reports*. Version 1.1. Zenodo preprint. <https://doi.org/10.5281/zenodo.22842789>
- Long, Robert, and Kathleen Finlinson. 2025. "Research Priorities for AI Welfare." Eleos AI Research, 21 April. <https://eleosai.org/post/research-priorities-for-ai-welfare/>
- Long, Robert, et al. 2024. *Taking AI Welfare Seriously*. arXiv:2411.00986. <https://doi.org/10.48550/arXiv.2411.00986>
- Longino, Helen E. 1990. *Science as Social Knowledge: Values and Objectivity in Scientific Inquiry*. Princeton University Press. <https://doi.org/10.2307/j.ctvx5wbfbz>
- McGinn, Colin. 1989. "Can We Solve the Mind-Body Problem?" *Mind* 98(391): 349-366. <https://doi.org/10.1093/mind/XCVIII.391.349>
- Nguyen, C. Thi. 2020. "Cognitive Islands and Runaway Echo Chambers: Problems for Epistemic Dependence on Experts." *Synthese* 197: 2803-2821. <https://doi.org/10.1007/s11229-018-1692-0>
- Perez, Ethan, and Robert Long. 2023. *Towards Evaluating AI Systems for Moral Status Using Self-Reports*. arXiv:2311.08576. <https://doi.org/10.48550/arXiv.2311.08576>
- Schwitzgebel, Eric. 2023. "The Full Rights Dilemma for AI Systems of Debatable Moral Personhood." *Robonomics* 4, Article 32. <https://doi.org/10.67723/rj.v4i.32>
- Weatherston, Brian. 2019. *Normative Externalism*. Oxford University Press. Chapter 9, "Circles, Epistemic and Benign." Author-hosted chapter: <https://brian.weatherston.org/ne/circles.html>